Precision–recall discrimination across forecasting horizons

Seven forecasting architectures | 18,839 aligned synthetic validation windows | Exact saved PR curves

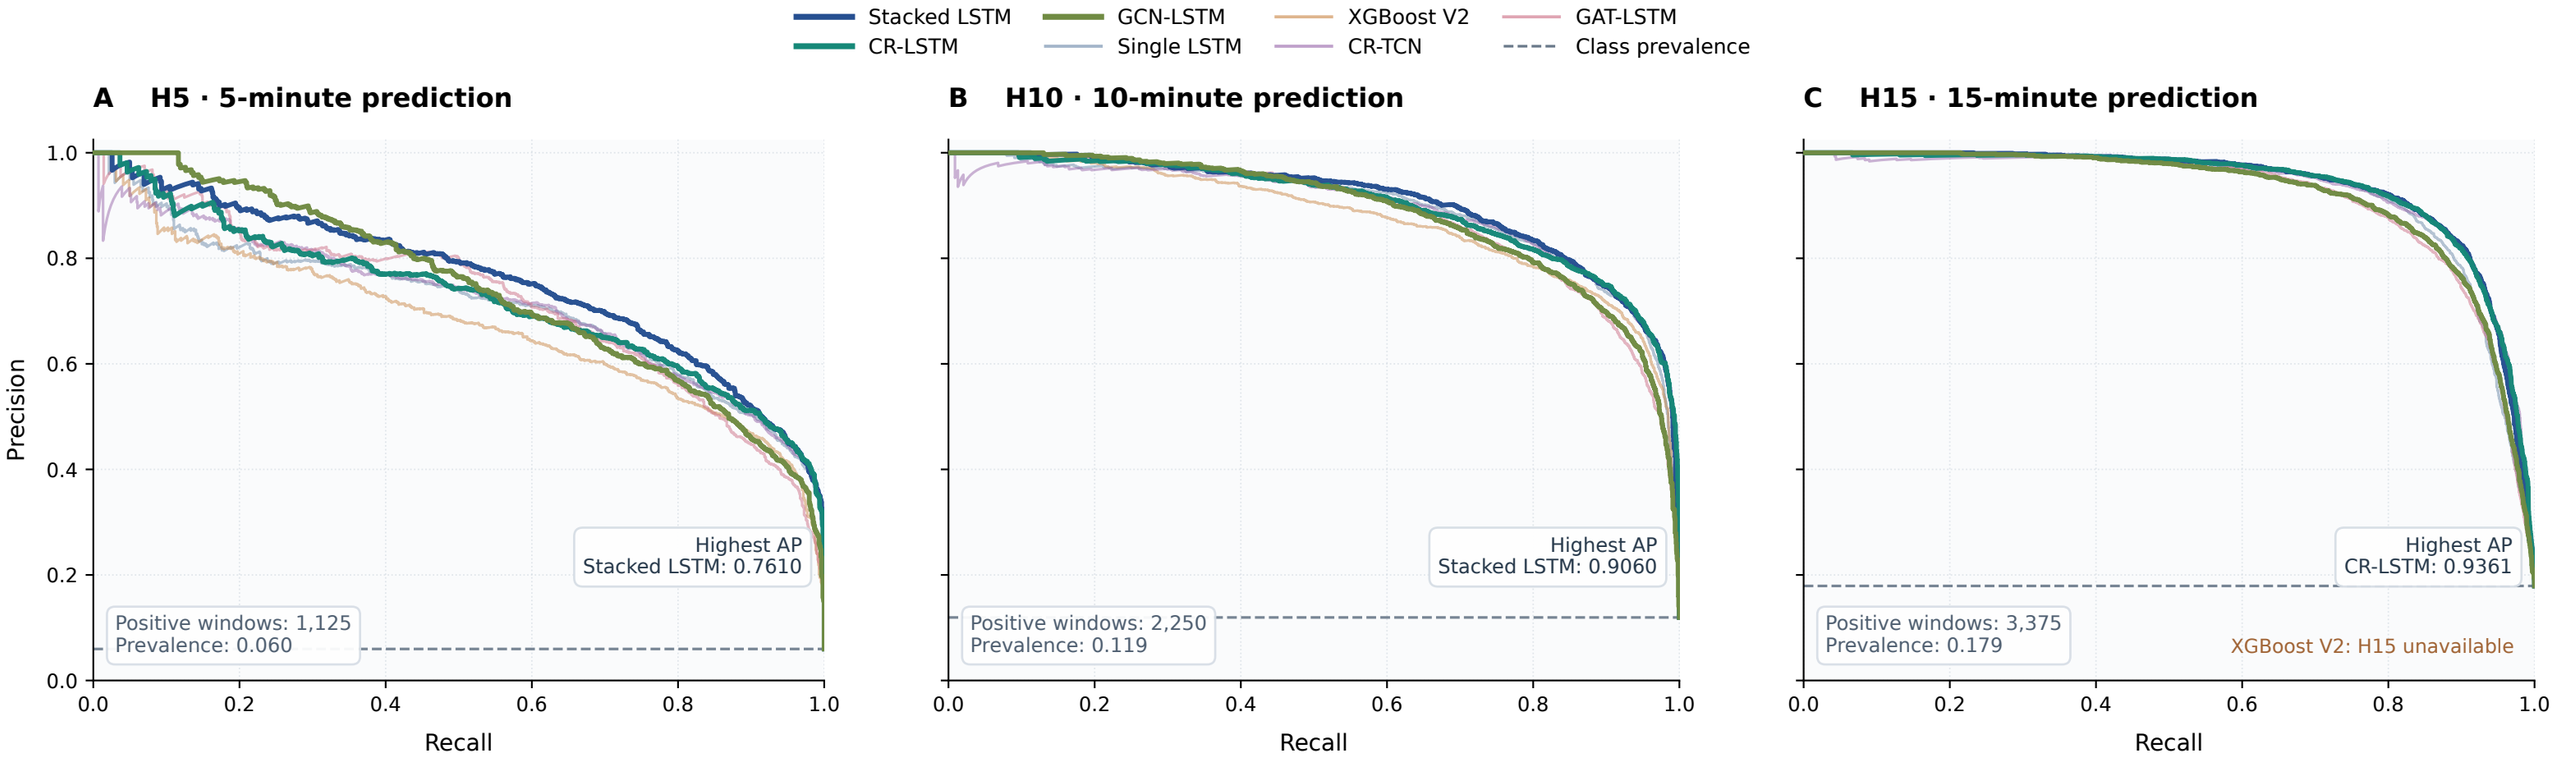


Figure 2. Precision–recall curves for 5-, 10-, and 15-minute future-attack prediction. Stacked LSTM, CR-LSTM and GCN-LSTM are emphasized; other evaluated architectures remain visible.

Average precision is computed from the exact saved curves and differs from overall classification accuracy. Positive prevalence increases with forecast horizon.

Synthetic NAF v2 development validation. The other Graph LSTM variants are compared in Table 1 but have no curves in the F09 archive. No independent-window confidence intervals are implied.