Source-dependent future attack-family classification

Conditional future-event classification | Two competing-risk models | Two synthetic validation source groups

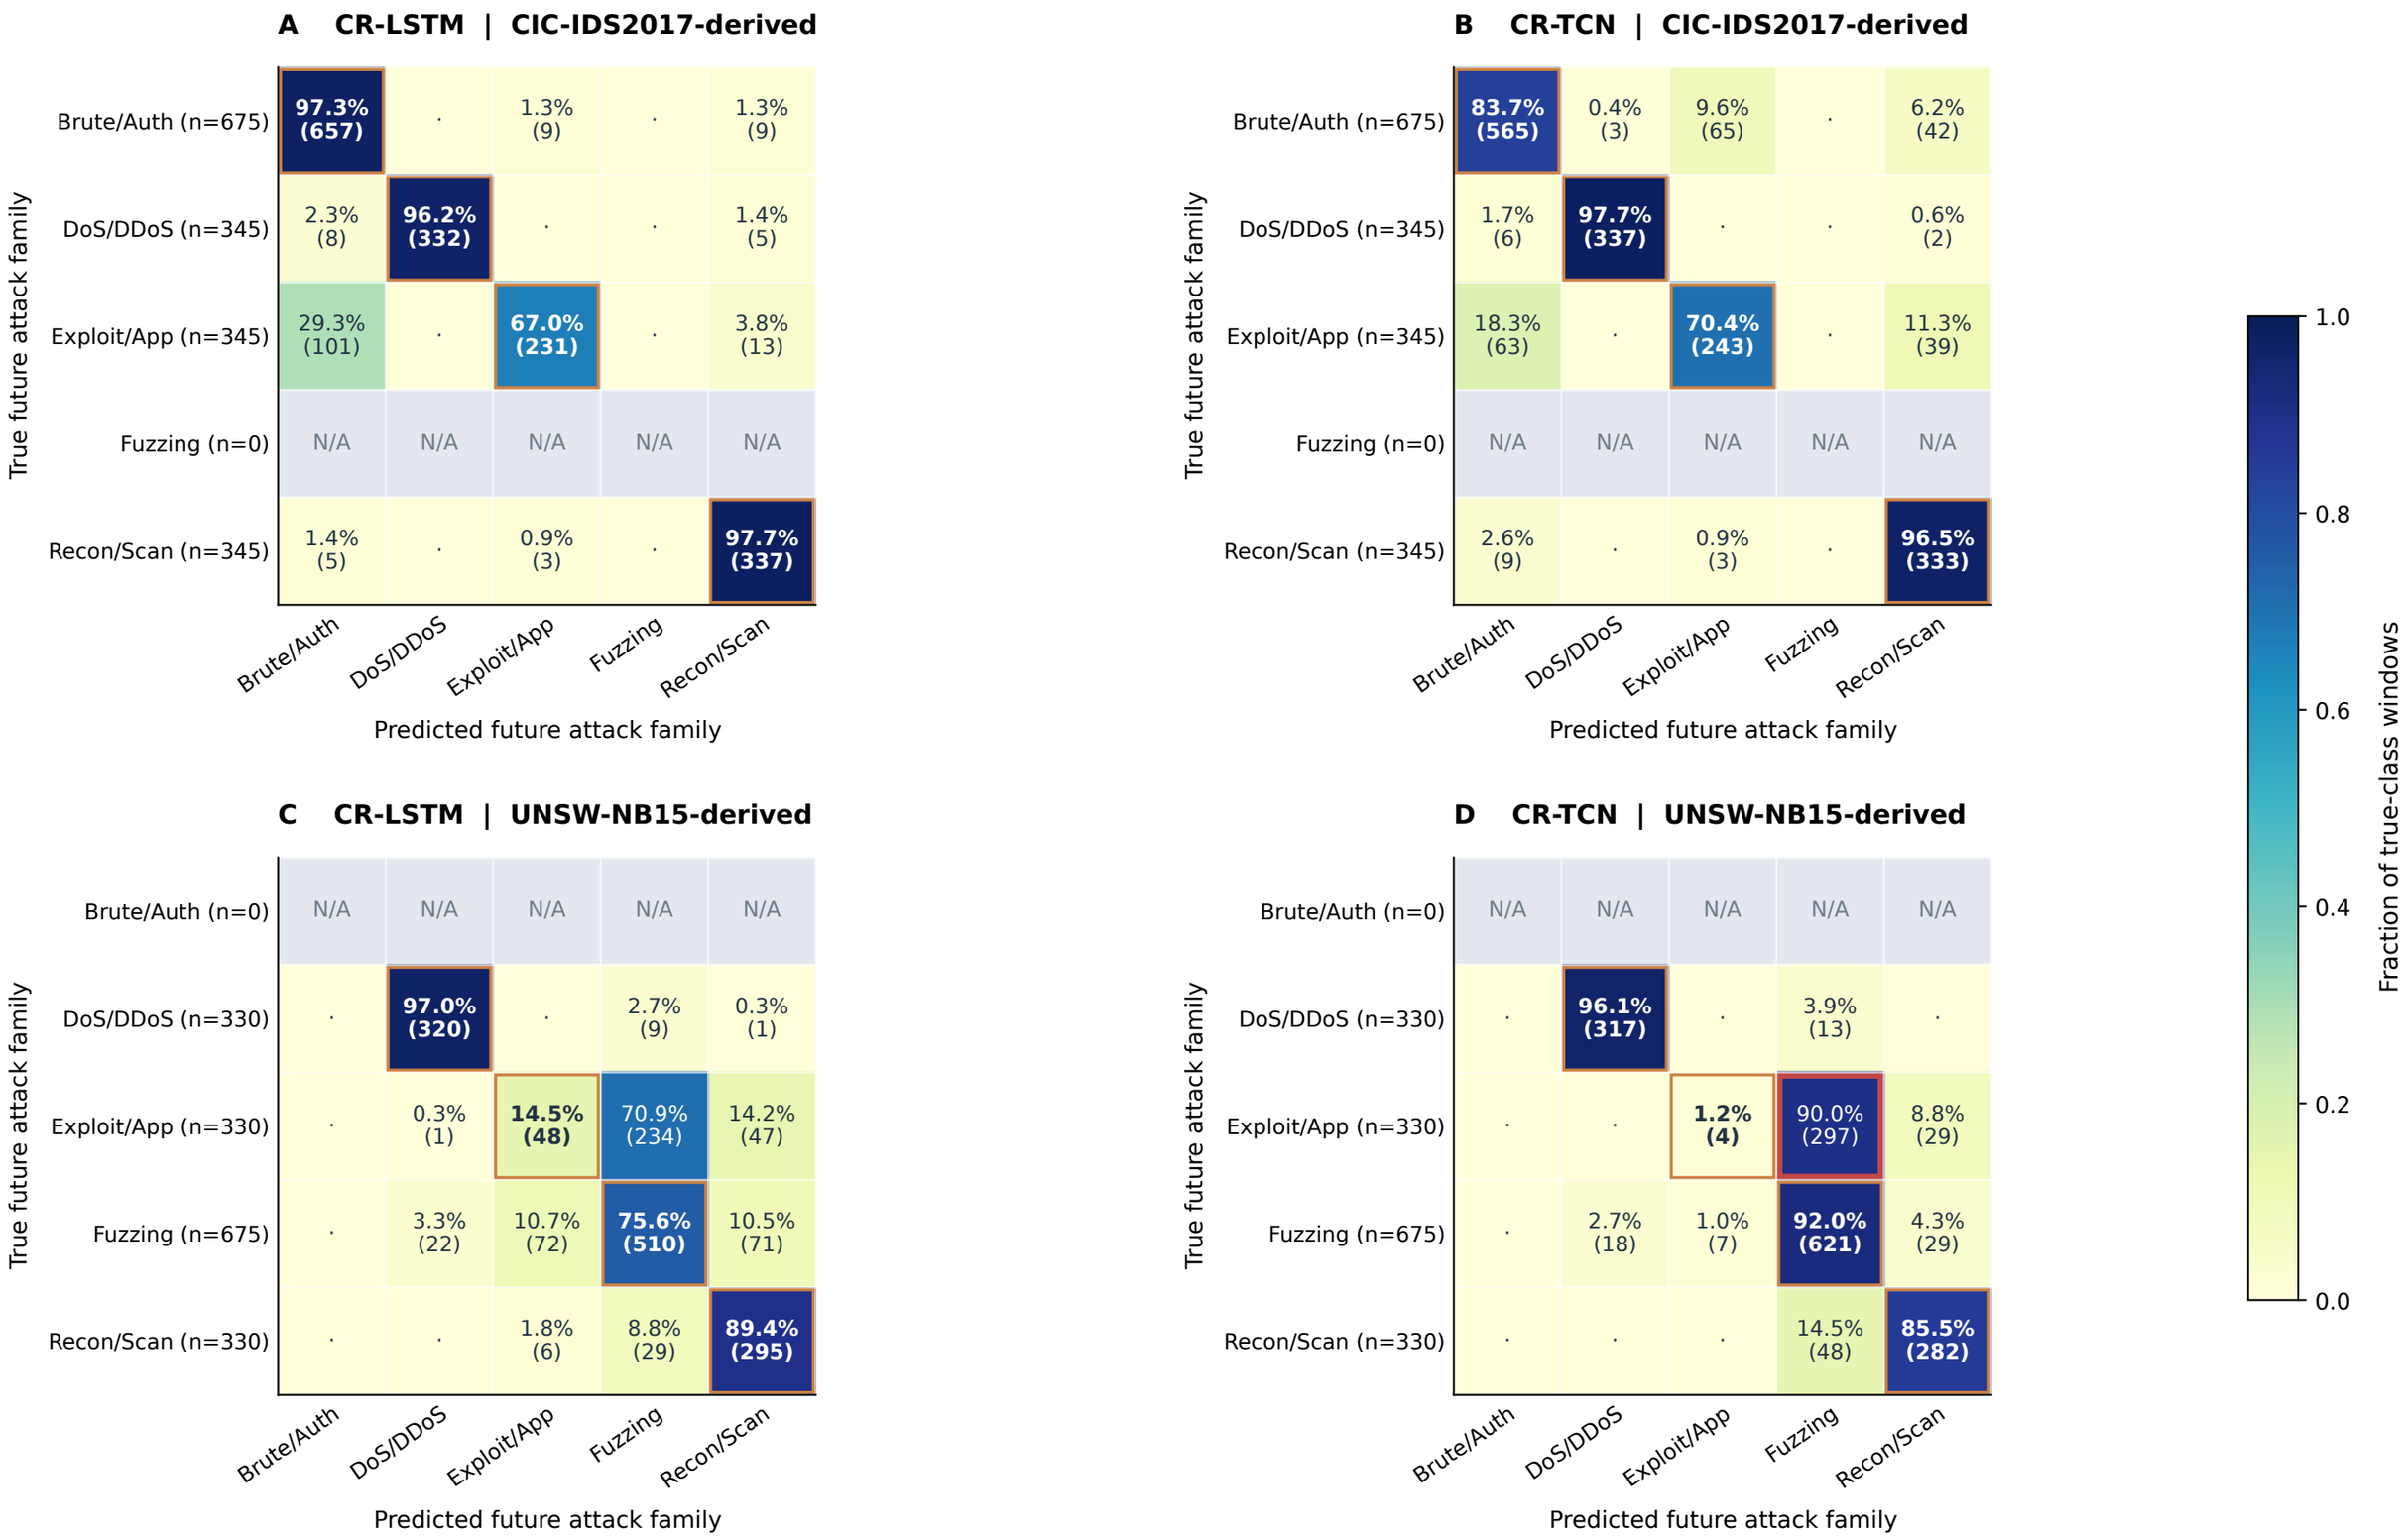

KEY FINDING On UNSW-NB15-derived Exploit/App windows, CR-LSTM correctly classified 48/330 (14.5%), versus 4/330 (1.2%) for CR-TCN.

CR-TCN assigned 297/330 (90.0%) of these windows to Fuzzing (red outlined cell). Orange diagonal outlines identify correct-family predictions.

Synthetic NAF v2 development validation. Rows are normalized by source-specific true-class support; gray rows denote absent classes, not zero recall. Family classification is conditional on a true future event.